



# OSN Tutorial 2

BUILDING YOUR OWN SHELL

# PROCESSES

- As we covered previous time, processes have their own unique ID with which they can be identified.
- Try typing “ps” on your terminal to find out the processes that are currently running and their pid.
- But the question is :Are all these processes running in a similar fashion?  
...  
...  
They are not.

# CONTROLLING TERMINAL

- To understand the difference between background and foreground processes, it is important for one to understand what a controlling terminal is.
- A controlling terminal is a terminal in Linux where a process starts. If you enter a command from the shell to “sleep”, the terminal window where this command is entered is the controlling terminal.
- Boiling it down, it's the terminal window that you are using.
- All input to the process is given through this terminal only, and all commands entered are run as processes in this terminal.
- Now, getting to foreground and background processes..

# FOREGROUND AND BACKGROUND

- Processes can be divided in two categories widely :
  - I. Foreground Processes
  - II. Background Processes
- By default, processes run in the foreground. To run any command in the background, type an ampersand (&) at the end of the command.
- A foreground process has **access to the controlling terminal**, and the shell waits for it to end before it can resume its operation.
- However, a background process as the name suggests runs in the background and has no user interaction and does not have access to the controlling terminal.
- It cannot read from stdin, but it CAN output to stdout.

# SWITCHING BETWEEN FG AND BG

- So.. a process switching from background process to a foreground process is basically gaining access to the controlling terminal.
- However, a background process cannot change the access of a controlling terminal normally or write to it.
- There is some signal handling that needs to happen here when a background process must access these functionalities.

# MORE PROCESSES

- We discussed `fork()` last time which creates a new process.
- Doing Mini Project 0, you might have noticed that the parent and child process go on to execute simultaneously after this call.
- However, having the parent wait before this child completes can be beneficial in some cases.
- So how do you make the parent wait for it's child process to end?

# MAKE IT WAIT


- `Wait()` is the system call that you are looking for here. This is a blocking command, and it makes the parent wait until one of its child terminates.
- `Waitpid(pid_t pid)` is another system call that can be used for this purpose which waits for the child with process id `pid` to terminate.
- `Waitpid` can take several more arguments and you should probably man it for further information (Hint : You will need it for Mini Project).

# REAL POTENTIAL OF FORK - EXECVP

Execvp replaces the current process image with a new process image that is specified in its arguments.

How is it useful?

Fork a new process from within a process. This process is duplicate of what you were running but then run Execvp and BOOM! You have a totally new process with a different functionality at hand.





# FILE DESCRIPTORS

- We discussed about it before. It's just a number that signifies a file.
- Some are already assigned : 0 (stdin), 1 (stdout) and 2 (stderr).
- Kernel maintains a table of open file descriptors for each process (which is also changed with context switch).
- This table maps this small integer to a file descriptor structure (struct fd in Linux) which has the required info about an open file.
- This structure also contains a pointer to an open file description,

# FILE DESCRIPTORS ALONG WITH FORK()

- One may ask this question, “What happens to a file descriptor after a `fork()` call?”
- While some may think that since the process created afterwards is an entirely different process from the original one, there will be no effect of one or the other.
- That’s wrong, however.
- An important detail of the design of Unix is that the open file description and the file buffer are kernel data structures.
- Therefore, these are NOT duplicated by a `fork()` call.
- Hence, whenever one of the processes reads from or writes to its local file descriptor, the effects in the open file description and file buffer are seen by the other process.

# FILES

- Now that we know about file descriptors, we can get to files. Fun fact : Everything in Linux is a file.
- There can be a file a.txt on your system, and it can have multiple open file description in the kernel which can be pointed by multiple file descriptors.
- It's important to understand that a file descriptor will only affect another file descriptor if they share the same open file description.
- By default, using open creates a new open file description.



Ólafur Waage  
@olafurw

Linux: "Everything is a file."

"Everything? Directories?"  
Linux: "File."

"Sockets? Devices?"  
Linux: "Yup, Files."

"My constant worry that I'll never  
be good enough?"  
Linux: "2 files actually."

Always has been

Wait it's all Files



# DUP SYSTEM CALL

- Basically, dupes a file descriptor that is passed as an argument.
- On success, returns a new file descriptor (-1 on error).
- It assigns the lowest available file descriptors.
- These two file descriptors can be used interchangeably, as they have the same open file description and share the same offset etc.
- Closing one does not affect the other.

## DUP2 SYSTEM CALL

- Dup2 system call is very similar in its functionality. It takes two arguments the oldfd and newfd and copies oldfd into newfd.
- If newfd was previously open, it closes it before reusing it, automatically.
- Dup2(oldfd, newfd)
- You can use these system calls for implementing I/O redirection.
- How? Figure it out.

# USING PIPES

- A pipe is a connection between two processes.
- Takes output of one and gives as input for another.
- Try it out yourself.
- This is on a crude level though; true piping goes a bit deeper than this.

## PIPE SYSTEM CALL


- Opens a pipe. Basically, a virtual file which can be written to and from.
- This pipe can be used by the process and its descendants since they get a copy of it.
- This system call basically returns two file descriptors, one pointing to the read end of the pipe and the other to the write-end.
- The read-end of the pipe can be used for reading from the pipe and write can be used to write to the pipe.



# SIGNAL()

- There is quite a bit of signals involved when multiple processes are trying to get their job done.
- These can be used to terminate a process, stop it etc.
- Using signal system call, you can change how a process reacts upon receiving a signal.
- You can create your own custom signal handler, choose to ignore the signal or choose to receive it.
- It's very useful here since you can ignore signals using this.

How signal handling works:

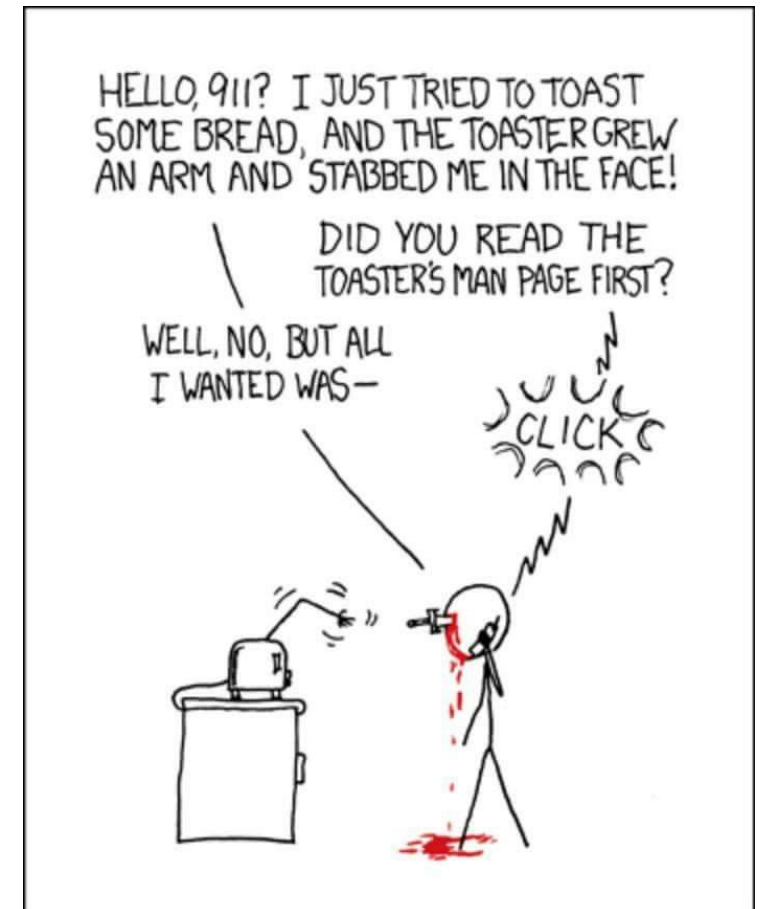


## SOME GENERAL ADVICE (SERIOUSLY THIS TIME)

Read MAN pages.

Write modular code.

Start on time.



Any Doubts?